

Lottery Sampling, Lottery Saving and other Algorithms for Top k Frequent Elements

Conrado Martínez Gonzalo Solera

September 9, 2018

Es muy buena idea que vayas escribiendo en este documento de trabajo u otro equivalente que constituya la base de tu documentación del TFG, las diferentes cosas que vayas haciendo, de la manera más organizada que puedas, aunque no esé muy pulido. Este documento te da un buen punto de partida sobre cómo organizar tu TFG (habrás de incluir algunas secciones “chorras”, you know it) y también va estableciendo el marco de trabajo (p.e. esquema de nombres para los diferentes algoritmos, la plantilla “unificada” para todos ellos, las definiciones de las medidas de rendimiento, ...)

1 Preliminaries

Consider a (large) data stream $\mathcal{Z} = z_1, \dots, z_N$, where each z_j is drawn from some domain or *universe* \mathcal{U} , and let $n \leq N$ be the number of distinct elements in \mathcal{Z} . We may thus look at the multiset X underlying \mathcal{Z}

$$X = \{x_1^{f_1}, \dots, x_n^{f_n}\}$$

where x_1, \dots, x_n are the n distinct elements that occur in \mathcal{Z} and f_i denotes the number of occurrences (absolute frequency) of x_i in \mathcal{Z} . We will assume, w.l.o.g., that we index the elements in X in non-increasing order of frequency, thus $f_1 \geq f_2 \geq \dots \geq f_{n-1} \geq f_n > 0$. We will use $p_i = f_i/N$ to denote the relative frequency of x_i . For simplicity, we will assume that $f_1 > f_2 > \dots > f_n$ in the definitions below—they can be more or less easily adapted to cope with elements of identical frequency.

The two problems that we want to study here are:

1. Top k most frequent elements. Given \mathcal{Z} and a value $k \leq n$, we want to find $\{x_1, \dots, x_k\}$ (or any subset of k distinct elements with maximal frequencies).
2. Heavy hitters. Given \mathcal{Z} and a value c , $0 < c < 1$, we want to find (or count) the number of distinct elements in \mathcal{Z} with relative frequency $p_i \geq c$.

Those elements are called *heavy hitters*. Given the data stream and the value c , we want to obtain $\{x_1, \dots, x_{k^*}\}$, where k^* is the largest value k such that $p_k \geq c$. The value k^* is the number of heavy hitters.

Moreover, in both problems, we might want that the algorithm returns the frequency f_i of the returned elements.

None of these two problems can be solved exactly unless we can keep $\Theta(n)$ elements in memory; thus under the tight memory constraints of the data stream model, we must aim at approximate good solutions. Hence, the algorithms that we describe next might return elements which are not among the most frequent elements, or that are not heavy hitters, and rather than the frequencies f_i of the returned elements, we will have to content ourselves with estimations f'_i of the real frequencies.

We will concentrate in algorithms for top k most frequent elements. Notice that there can be at most $\lceil 1/c \rceil$ heavy hitters in a data stream, and thus an algorithm that retrieves the top $k^* = \lceil 1/c \rceil$ most frequent elements will obtain all the heavy hitters.

2 The Algorithms

The algorithms that we shall consider next fail under the following scheme. They all keep a *sample* \mathcal{S} of up to m distinct elements, for some value $m \geq k$ fixed in advance. Each element x in the sample is equipped with a frequency counter

$$f'(x) = f_{\text{obs}}(x) + f_{\text{ini}}(x),$$

which is the sum of two components: $f_{\text{obs}}(x)$ is the number of times that x has been observed in the data stream since it has been included in \mathcal{S} ; $f_{\text{ini}}(x)$ is an estimation of the number of occurrences of x prior to the occurrence in which x has been added to the sample.

In the initial phase, the algorithms populate the sample with the initial distinct elements in the data stream, recording their frequencies. In particular, $f_{\text{ini}}(x) = 0$ for these elements, and $f'(x)$ will be the real frequency of x in the “prefix” of the data stream examined so far. This first phase ends when m distinct elements have already been collected and an $(m + 1)$ -th distinct element occurs in the data stream, or when the data stream is exhausted—because $m \geq n$. In the latter case, the sample has complete information about the full data stream and all required information can be obtained exactly. In the interesting case, when $m < n$ (typically, $m \ll n$), the algorithms loop through the remaining items in the data stream \mathcal{Z} . For every incoming item z , if z is already in \mathcal{S} , the algorithms update $f'(z)$ and perhaps perform some additional bookkeeping (depending on the particular algorithm under consideration), but not much more has to be done. On the other hand, if $z \notin \mathcal{S}$ then we apply some criterium $\text{add?}(z, \mathcal{S})$; if $\text{add?}(z, \mathcal{S})$ is not satisfied, then z is discarded. Otherwise, some element $z' \in \mathcal{S}$ is chosen and evicted from the sample ($z' = \text{element_to_evict}(\mathcal{S})$), and z is added to \mathcal{S} ; the frequency counter is initialized

$f'(z) = 1 + f_{\text{ini}}(z)$. The algorithms that we will consider differ hence in three main aspects:

1. When do we add to the sample an incoming item $z \notin \mathcal{S}$ (**add?**(z, \mathcal{S})? In several algorithms that we will consider later, this is a randomized choice.
2. Which element z^* do we evict from the sample when a new element z has to be added to the sample (**element_to_evict**(\mathcal{S}))? Again, this choice might be probabilistic.
3. How do we estimate the frequency $f_{\text{ini}}(z)$, that is, the number of prior occurrences of a new element z that is to be added to the sample?

Take for instance the well-known SPACESAVING (SS) [?]. It always adds incoming items $z \notin \mathcal{S}$, that is, **add?**(z, \mathcal{S}) = **true**. The element z^* selected for eviction is one with lowest estimated frequency

$$z^* = z_{\text{lef}} = \arg \min_{x \in \mathcal{S}} \{f'(x)\}.$$

And the newly added item z inherits the frequency of the evicted item: $f_{\text{ini}}(z) = f'(z^*)$.

Let us now consider our first new algorithm LOTTERYSAMPLING (LS). Each element x in the sample has, besides its frequency counter, a ticket $t(x) \in [0, 1]$. When a new item z in the data stream is retrieved LS generates, uniformly at random in $(0, 1)$, a ticket t for z . If z was already in the sample its frequency counter $f'(z)$ is updated (as usual), but also its ticket: if $t > t(z)$ then $t(z) := t$. If z is not in the sample, we compare t with the minimum ticket $t_{\min} = t(z_{\min})$ in the sample. If $t > t_{\min}$ then z is added to the sample, otherwise z is discarded. The element z^* to be evicted is z_{\min} , the one with the minimum ticket, and the frequency counter of z is initialized with

$$f_{\text{ini}}(z) = \left\lfloor \frac{1}{1 - t_{\min}} \right\rfloor.$$

LOTTERYSAVING-LOWESTESTIMATEDFREQUENCY (LS-LEF) combines some of the ideas of the two algorithms above. When an item z is not in the sample we compare its ticket t with the ticket $t_{\text{lef}} = t(z_{\text{lef}})$ of an element z_{lef} with smallest frequency counter (in case of ties, any of them is picked at random). Notice that t_{lef} might be or not the minimum ticket in the sample. If $t > t_{\text{lef}}$ then z is added and $z^* = z_{\text{lef}}$ is evicted, and $f_{\text{ini}}(z) = f'(z^*)$.

LOTTERYSAVING-LEASTRECENTLYOBSERVED (LS-LRO) is similar to LS-LEF, but the ticket t a candidate element z to enter the sample is compared with the ticket $t_{\text{lro}} = t(z_{\text{lro}})$ of the item z_{lro} in the sample that was observed least recently (i.e., z_{lro} is the item with frequency counter updated the longest ago). Like in LS-LEF, if $t > t_{\text{lro}}$ then z is added and $z^* = z_{\text{lro}}$ is evicted. The counter of z is initialized using $f_{\text{ini}}(z) = f'(z_{\text{lro}})$.

In LOTTERYSAVING-SMALLESTTICKET (LS-ST) the element to be evicted is z_{\min} ; but instead of the rule for initialization of f' of LS it uses the estimated

frequency f' of z^* , like LS-LEF and LS-LRO. That is, when an item z is not in the sample we compare its ticket t with the smallest ticket $t_{\min} = t(z_{\min})$ of the element z_{\min} with smallest tickets (w.l.o.g. we can safely assume that no two tickets ever are identical). If $t > t_{\min}$ then z is added and $z^* = z_{\min}$ is evicted, like in LS. But, unlike LS, we have $f_{\text{ini}}(z) = f'(z^*)$.

LOTTERYSAVING-THRESHOLD (LS-THR- θ) has a parameter $\theta \in [0, 1]$. We omit the parameter θ in the name, thus write LS-THR in generic discussions about this algorithm. A new item $z \notin \mathcal{S}$ is added to the sample if and only if the ticket t generated for z is larger than $\theta \cdot t_{\max}$, where t_{\max} is the largest ticket among the elements in the sample (LS-THR-0 is equivalent to SS).

If z is added to \mathcal{S} , LS-THR evicts the element $z^* = z_{\text{lef}}$ with minimum frequency counter and z inherits the frequency counter of z^* as the initial estimation: $f_{\text{ini}}(z) = f'(z^*)$.

LOTTERYSAVING-ABOVEMEAN (LS-AM) adds a new item to the sample if its ticket t is larger than the mean value \bar{t} of all the tickets in the sample. When $t > \bar{t}$, the evicted item z^* is z_{lef} , one with smallest frequency counter, and z inherits $f'(z^*)$ like in LS-LEF and LS-THR.

Another variation is LOTTERYSAVING-ABOVEMEDIAN (LS-MED) which adds a new item to the sample if its ticket t is above the median t_{med} of the tickets in the sample. The remaining choices for the algorithm are as in LS-AM. LS-MED can be generalized to consider the α -quantile (LS-QUANT- α) of the tickets in the sample. Notice that LS-MED is LS-QUANT-0.5; the algorithm uses $t_{(\lceil \alpha \cdot m \rceil)}$ to decide whether to include or not a new item in the sample, where $t_{(r)}$ denotes the r -th smallest ticket in the sample, $1 \leq r \leq m$.

The basic variants of LS-LEF and LS-LRO can be modified to incorporate the ideas behind LS-AM and LS-MED. Thus LS-LEF-AM and LS-LRO-AM will add an element $z \notin \mathcal{S}$ to the sample if $t > t(z^*)$ **or** t is above the mean value \bar{t} of the tickets in the sample. The evicted item z^* and the initialization of the frequency counter of z is as in the corresponding basic algorithms.

Instead of the mean value \bar{t} of the tickets in the sample, we might consider the median of the tickets, or more generally, the α -quantile of the tickets, getting the variants LS-LEF-MED LS-LRO-MED, or more generally, LS-LEF-QUANT- α and LS-LRO-QUANT- α .

Notice that it makes no sense to “combine” LS-AM and LS-QUANT with LS-ST since $t_{\min} < \bar{t}$ and $t_{\min} \leq t_{(r)}$ for all r .

Another source of variations stems from the way tickets are updated when an item z already in the sample is the incoming element from the stream. In all cases above we have assumed that tickets are updated via $t(z) := \max(t, t(z))$, where t is the ticket generated for the current instance of z , whereas $t(z)$ is the ticket associated to z in the sample. We can derive corresponding variants of all the others discussed above where the items retain the ticket t with which they entered the sample, and never update it once in the sample, for instance.

All algorithms using tickets can be also characterized as defining a *threshold value* t^* and a candidate item z^* for eviction. The criterium for addition of an element $z \notin \mathcal{S}$ is always whether $t > t^*$ or not, where t is the ticket that was generated for the current instance. For example, lottery sampling (LS) takes

Algorithm	t^*	z^*	$f_{ini}(z)$
SS	0	z_{lef}	$f'(z^*)$
LS	$t_{\min} = t(z^*)$	z_{\min}	$\lfloor \frac{1}{1-t^*} \rfloor$
LS-LEF	$t(z^*)$	z_{lef}	$f'(z^*)$
LS-LRO	$t(z^*)$	z_{lro}	$f'(z^*)$
LS-ST	$t(z^*)$	z_{\min}	$f'(z^*)$
LS-THR	$\theta \cdot t_{(m)} = \theta \cdot t_{\max}$	z_{lef}	$f'(z^*)$
LS-AM	$\bar{t} = \sum_{z \in \mathcal{S}} t(z)/m$	z_{lef}	$f'(z^*)$
LS-QUANT	$t_{\lceil \alpha m \rceil}$	z_{lef}	$f'(z^*)$
LS-LEF-AM	$\min\{\bar{t}, t(z^*)\}$	z_{lef}	$f'(z^*)$
LS-LEF-MED	$\min\{t_{\lfloor (m+1)/2 \rfloor}, t(z^*)\}$	z_{lef}	$f'(z^*)$

$t^* \equiv t_{\min} = \min\{t(z) \mid z \in \mathcal{S}\}$, and z^* an element with the minimum ticket t_{\min} . Recall that we can safely assume that all tickets in \mathcal{S} are distinct.

With these conventions the template for all the algorithms discussed above is as follows:

```

Populate S with the first m distinct elements in Z,
    updating tickets and frequencies
while Z is not exhausted do
    z := current item in Z
    t := Random(0,1)
    if z in S then
        Update the ticket t(z)
        f'(z) := f'(z) + 1
    else if t > t* then
        S := S - {z*} + {z}
        t(z) := t
        f'(z) := f'(z) + 1
    else // do nothing
        Update t* and z*
endwhile
Report the k elements in S with largest f'

```

Table ?? summarizes the characteristic values of t^* and z^* . Observe that we have considered two options for the initialization of $f_{ini}(z)$ and three options for z^* (z_{\min} = the element with minimum t , z_{lef} = an element with minimum f' , and z_{lro} = the least recently observed element); if we contemplate C different ways to define t^* ($C = 9$ in our table), we could consider then up to $6C$ different possible combinations—some make no sense. If we add on top of that the choice to update or not tickets of sampled elements we raise the count to $12C$. Moreover, LS-THR and LS-QUANT depend on an additional real parameter in $[0, 1]$ giving us an additional source of variability.

The three strategies LS-ST, LS-LEF and LS-LRO are almost identical, they only differ in the definition of the element z^* to be evicted. These three strategies will be called *canonical*, as all the other algorithms that we have discussed here

can be seen as variations of one of the three; maybe in the way that tickets are updated, maybe in the way that t^* is defined (e.g., $t^* \neq t(z^*)$), or the way that f' is initialized (e.g., $f_{ini}(z) \neq f'(z^*)$).

It is clear that the combinations that we can make “playing” around with the criteria for addition ($\text{add?}(z, \mathcal{S})$), the criteria for eviction and the initialization of the frequency counter are almost endless, but here we will restrict ourselves to the those that we have been able to analyze and for which the experiments give the best results. Another important issue in our choice has been the efficiency with which the different operations can be supported. Thus for instance, to implement LS-LEF it is very convenient to maintain the elements of \mathcal{S} sorted by frequency counter; locating the least frequent element becomes trivial, and maintaining the order of the sample after each frequency update is also very easy and efficient, as they increment one by one.

Gonzalo: agrega a esta seccion otros algoritmos que estes probando. En el documento final es posible que hayas de hacer una “purga”, eliminando las opciones que no han dado buenos resultados. Sin embargo, SpaceSaving, Lottery Sampling y las tres variantes canonicas LS-ST, LS-LEF y LS-LRO deben estar siempre, como referencia. Para nuevas variantes inventa nombres y acronimos. Tambien habriamos de pensar algun modo facil de indicar si una estrategia es con o sin update de tickets, en vez de inventar nuevos nombres. Por ejemplo, LS*-LEF es la variante de LS-LEF sin hacer update de los tickets. Por cierto que nn una primera impresion yo diria que LS*-LEF es casi lo mismo que SpaceSaving pero lanzando una moneda e incluyendo al item z nuevo en el sample con probabilidad $m/(m+1)$ (con probabilidad $1/(m+1)$ el item nuevo no reemplaza a z_{lef}).

2.1 Implementing the Algorithms

3 Measures of Quality

Of course, for a top k frequent elements algorithm we would like that it reports the true k most frequent elements and their respective frequencies, but we know that this will not be possible unless we had a linear amount of memory ($\Theta(n)$)—this includes the trivial situation when $m \geq n$. Hence we will assume in what follows that $m \ll n$, and we would like that our algorithms return good approximations: elements that are the most frequent or close to that, and good estimations of their respective frequencies. To measure the quality of the different algorithm we introduce several measures, in which two random variables will be essential:

1. The indicator variable Y'_i tells us whether x_i has been sampled or not: $Y'_i = 1$ if $x_i \in \mathcal{S}$ and $Y'_i = 0$, otherwise. Notice that the Y'_i 's are not independent and $Y'_1 + \dots + Y'_n = m$.

2. Since $m \geq k$, the algorithm has to choose k out of the m elements to report them as the top k most frequent elements. Quite naturally, in all our algorithms the k elements with largest frequency counter will be the ones to be reported and their estimated frequencies will be those collected by the algorithm. The indicator variable $Y_i = 1$ if x_i is reported, $Y_i = 0$ otherwise (either because $x_i \notin \mathcal{S}$ or because $f'(x_i)$ is not among the largest k values of f' in \mathcal{S}). Like the Y_i' 's, the Y_i 's are not independent; and $Y_1 + \dots + Y_n = k$.

Our first measure, (*weighted*) *recall*, measures the proportion of top frequent elements actually found by the algorithm:

$$R = \frac{p_1 Y_1 + \dots + p_k Y_k}{p_1 + \dots + p_k}$$

A proxy for R is given by

$$R' = \frac{p_1 Y_1' + \dots + p_k Y_k'}{p_1 + \dots + p_k},$$

which is always an upper bound for R , since $Y_i \leq Y_i'$ for all i , $1 \leq i \leq n$. Notice that, in all cases, $0 \leq R \leq 1$: $R = 0$ if no true frequent element is reported, and $R = 1$ if exactly the top k most frequent elements are reported. The weights p_i will penalize more algorithms that miss the most frequent elements among the top frequent, and will lessen the impact in R if the algorithm missed less frequent elements (albeit among the most frequent). The same is true for R' .

Another measure is (*weighted*) *precision*, which gives us an indication of the quality of the reported elements:

$$P = \frac{p_1 Y_1 + \dots + p_n Y_n}{p_1 + \dots + p_k}$$

The numerator sums the weights of all reported elements; since exactly k elements are reported $0 < P \leq 1$. The proxy P' for P needs a different “normalizing” denominator:

$$P' = \frac{p_1 Y_1' + \dots + p_n Y_n'}{p_1 + \dots + p_m}$$

We can also use the “ordinary” measures of unweighted recall R_u and unweighted precision P_u used in most of the literature (e.g. in the experimental study of SS []). In terms of our indicator random variables we have

$$R_u = \frac{Y_1 + \dots + Y_k}{k}$$

$$P_u = \frac{Y_1 + \dots + Y_k}{m},$$

and the approximations (actually, upper bounds)

$$R'_u = \frac{Y_1' + \dots + Y_k'}{k}$$

$$P'_u = \frac{Y_1' + \dots + Y_k'}{m}.$$

Besides the measures of “recall” and “precision” we also consider measures for the quality of the estimated frequencies.

Error type I ϵ_I measures the deviation of the estimated frequencies for the frequent elements:

$$\epsilon_I = \frac{\sum_{1 \leq i \leq k} (f_i - \hat{f}_i)^2}{\sum_{1 \leq i \leq k} f_i^2},$$

where $\hat{f}_i = f'(x_i)$ if $x_i \in \mathcal{S}$ and $\hat{f}_i = 0$, otherwise. For reasons that will be discussed later we might also take $\hat{f}_i = \min\{f'(x_i), 2f_{\text{obs}}(x_i)\}$. Since $f_{\text{obs}}(x_i) \leq f_i$ for all i , the numerator in ϵ_I cannot exceed the denominator and thus ϵ_I will range between 0 (perfect estimation and all frequent elements in the sample) and 1 (all frequent elements missing and/or grossly overestimated). This measure does not take into account what are the non-frequent elements that get sampled or how good or bad are the estimates of their frequencies. For that we have a second measure, error type II. In this case the absolute errors in the estimates are not that useful; we should penalize more severely the errors in the estimation of elements of very low frequency since that might imply reporting them as frequent element and they are not even close. On the other hand, no penalty should be added if the element is not sample. Hence for the definition of error type II we will be doing better using relative errors, and weighting respect their presence or not in the sample:

$$\epsilon_{II} = \frac{1}{m} \sum_{i \geq k} Y'_i \left(\frac{\hat{f}_i - f_i}{f_i} \right)^2.$$

Since each term in the sum above is at most 1 and there are at most m non-null terms, ϵ_{II} also ranges between 0 (no mistakes in the frequency estimation) and 1 (all sampled items are infrequent and we make a gross overestimation of their frequencies).

Revisa las definiciones para ver si son correctas y se adecuan con las que usas en los experimentos. Agrega nuevas medidas que estes usando y se te ocurran. Ya habra ocasion de “purgar” aquellas que resulten poco relevantes, redundantes o demasiado ad-hoc (como comentamos tiene que haber un fundamento racional que justifique la definicion de la medida, no vale que de los resultados que nos gustan :-))

R_u y P_u las versiones no ponderadas de *recall* y *precision* no son las mas utiles, pero han de estar y aparecer los valores correspondientes en la parte experimental. En el paper de SS, son las medidas que se usan y estara bien hacer la comparativa en base a estas dos medidas tambien.

4 Analysis of the Algorithms

5 Experimental Results

Esta sección (que ocupará una parte importante de este documento y también en tu TFG) se dividirá en no menos de tres partes:

- Descripción de los conjuntos de datos reales y sintéticos usados en los experimentos
- Descripción del *setup* experimental: cuantas veces se repetía el experimento, entorno de trabajo, etc.
- Presentación (tablas, gráficas) y análisis de los resultados experimentales sobre calidad de los algoritmos (recall, precision, errores de estimación de frecuencias, ...)
- Presentación (tablas, gráficas) y análisis de los resultados experimentales sobre rendimiento (tiempo de proceso, memoria, ...)